

# CoT or ReAct: Evaluating ReAct-Based LLM Reasoning for Korean Date and Schedule Processing

Yein Oh, Jisu Kim, Jihyun Jung, Sooyeon Park



## Table of Contents

- I. Introduction
- II. Tasks
- III. Methodology
- IV. Experiments
- V. Results
- VI. Conclusion



# I. Introduction

# Research Background – Korean Date and Schedule Processing

- NLP-based **schedule calculation** and **time normalization** are core functions in various AI services.
- Korean presents a high degree of difficulty due to complex constraints, including
  - **relative tense**
  - **complex combination of public holidays, weekdays, and intervals**
- Existing LLM-based approaches using **few-shot learning**, **Chain-of-Thought (CoT) reasoning**, and **fixed JSON output schemas** have achieved a moderate level of performance.
- However, they have shown **limitations** in consistently and accurately satisfying these constraints.



# Research Background – Korean Date and Schedule Processing

year		month		day		week		time		
Rules	Value	Rules	Value	Rules	Value	Rules	Value	Attribute	Rules	Value
N년(앞 전)	year="-N"	N개월(앞 전)	month="-N"	N일(앞 전)	day="-N"	N주(앞 전)	week="-N"	hour	N시간(앞 전)	hour="-N"
N년(뒤 후)	year="+N"	N개월(뒤 후)	month="+N"	N일(뒤 후)	day="+N"	N주(뒤 후)	week="+N"		N시간(뒤 후)	hour="+N"
내년 다음 (해 년)	year="+1"	(내 다음)달	month="+1"	내일 다음날	day="+1"	(저번 이번  다음)주	week="-1", "0", "+1"		N시간째	hour="~N"
작년 (저번  지난 이전)해	year="-1"	(저번 지난  이전)달	month="-1"	어제 이전날	day="-1"	N주(차 째)	week="~N"	minute	N분(앞 전)	minute="-N"
내후년	year="+2"	이번달	month="0"	(이틀 사흘 ... 뒤)	day="+2", "+3", ...				N분(뒤 후)	minute="+N"
재작년	year="-2"	N개월(차 째)	month="~N"	(이틀 사흘 ... 전)	day="-2", "-3", ...				N분째	minute="~N"
이번 (해 년) 올해	year="0"			모레 (사흘  글피) 나흘	day="+2", "+3", "+4"	second	N초(앞 전)	second="-N"		
N년(차 째)	year="~N"			그저께	day="-2"		N초(뒤 후)	second="+N"		
				N일(차 째)	day="~N"		N초째	second="~N"		



# Research Background – Korean Date and Schedule Processing

- However, they have shown **limitations** in consistently and accurately satisfying these constraints.

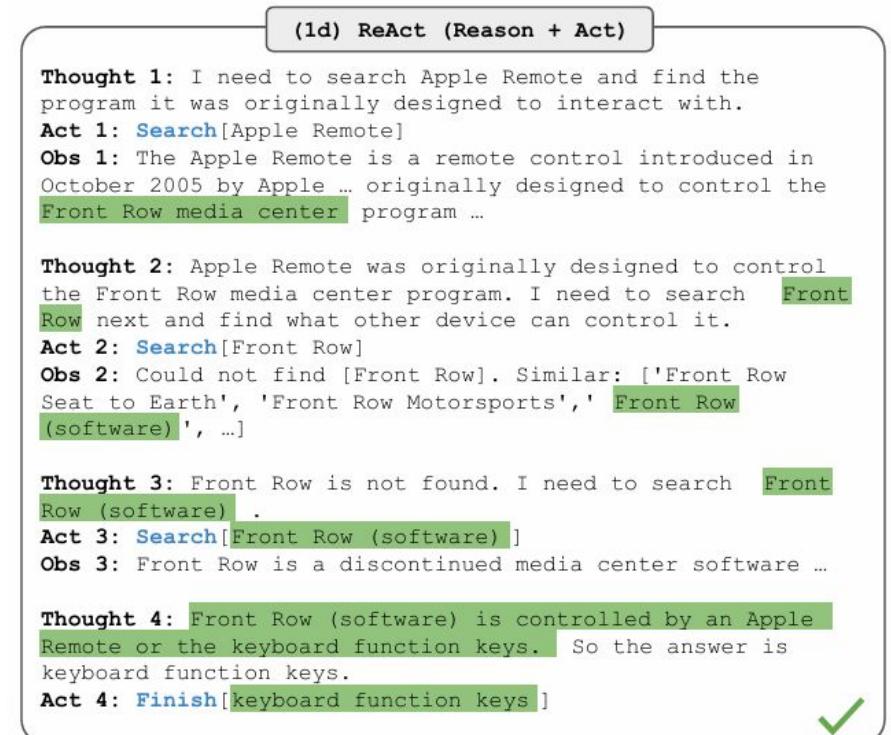
Attribute	Accuracy(%)
year	73.14%
month	61.21%
day	74.57%
week	91.39%
hour	75.00%
minute	50.00%
second	100.00%



EWHAWOMANS UNIVERSITY

# Proposal – ReAct

- ReAct = **Reason** + **Act**
  - Loop structure of “**Reasoning**” ↔ “**Action**” ↔ “**Observation**”
  - Incorporates external information through tool calls
    - Enables more structured and logical reasoning compared to simple CoT prompting



## II. Tasks

# Three Tasks for Evaluating Korean Temporal Reasoning

- **Constraint-based Korean Schedule Generation**
  - : to accurately **interpret complex temporal and scheduling constraints** within a **Korean** natural language query, in order to generate a list of dates that satisfies all specified conditions
    - Input: "수능 다음날부터 공휴일 제외하고 3일 간격으로 2개 날짜 제안해줘"
    - Output: a **list of dates** that **satisfy the given constraints**
- **Three Tasks (Increasing in Complexity)**
  - **T1**: Convert relative time expressions, **phrases** → absolute dates
  - **T2**: Convert relative time expressions, grounded in context, using complete **sentences** → absolute dates
  - **T3**: Generate **schedules** satisfying **multiple interacting constraints**



# Task Details

To compare CoT vs. ReAct robustness, we design **three tasks**, each isolating a different reasoning challenge:

- **T1 — Date Normalization (Phrases)**
  - Input: **standalone relative expression**
    - e.g., “이번 달 마지막 수요일”
  - Requires sequential temporal arithmetic (month-end → next weekday).
  - Output: **one absolute date** (YYYY-MM-DD).
  - Tests **core temporal arithmetic** without linguistic ambiguity.
- **T2 — Date Normalization Using Sentences**
  - Input: **full natural-language sentences** containing temporal phrases.
    - e.g., “2주 뒤인데, 만약 그날이 공휴일이면 다음 평일로 해줘.”
  - Requires finding the relevant span **before** calculating the date.
  - Same output format as T1.
  - Tests **linguistic grounding + disambiguation**.
- **T3 — Constraint-based Schedule Generation**
  - Input: **multiple explicit constraints** (intervals, weekdays, exclusions).
    - e.g. “내년 1월 첫 번째 금요일부터 시작해서 5일 간격으로 4회 일정 잡아주세요”
  - Must iteratively generate dates and filter invalid ones.
  - Output: **list of dates**.
  - Tests **long-horizon reasoning, constraint satisfaction, rule integration**.



# Hypotheses

- **H1:** The more **complex** the constraints, the **higher** the **accuracy** and constraints satisfaction in **ReAct** Agent.
- **H2:** **ReAct** may **increase response time** and **token usage**, but shows **Pareto-efficient** accuracy-to-cost performance.



### III. Methodology

# CoT: Chain-of-Thought Reasoning

Few-shot Chain-of-Thought (CoT) Prompting	Purely Internal Reasoning	Task-specific Formatting	Purpose as a Baseline
<ul style="list-style-type: none"><li>• Each task uses a small set of demonstrations showing the desired reasoning format.</li><li>• Model is instructed to produce step-by-step reasoning followed by a final answer.</li></ul>	<ul style="list-style-type: none"><li>• All computations occur inside the LLM.</li><li>• No external tools, symbolic modules, calculators, or validation steps.</li><li>• Represents the natural behavior of text-only reasoning.</li></ul>	<ul style="list-style-type: none"><li>• T1 &amp; T2: Output is a <b>single normalized date</b> (YYYY-MM-DD).</li><li>• T3: Output is a <b>sequence of dates</b> satisfying constraints.</li><li>• Prompts differ only in the input/output description; core CoT style is consistent.</li></ul>	<ul style="list-style-type: none"><li>• Provides a clean comparison against ReAct.</li><li>• Highlights how agentic stepwise execution can improve temporal reasoning beyond static CoT.</li></ul>



# ReAct Model: Thought → Action → Observation Loop

## Core Mechanism

- Alternates between:  
**Thought** (internal reasoning) → **Action** (tool call) → **Observation** (tool output).
- Modular design: planner (thought), tool executor (action), decision module (observation evaluation).
- Same loop for all tasks, but behavior differs by task complexity.

## Tools Used Across Tasks

- **Calculator** — date arithmetic (offsets, intervals, weekday shifts).
- **Calendar\_db** — holidays, lunar/solar terms, anniversaries.
- **Search** — event lookup when rules are non-deterministic.
- Modular separation supports multi-step reasoning in T3 while keeping T1/T2 lightweight.

### T1: Phrase-Level Date Normalization

- Input: isolated temporal phrase.
- Single-step reasoning → one tool call (e.g., calculator, calendar\_db).
- No iteration needed; observation directly yields normalized date.
- Focus: pure temporal arithmetic.

### T2: Sentence-Level Date Normalization

- Must first locate the temporal span inside a full sentence.
- Thought module extracts the temporal phrase and reformulates it into canonical tool input.
- One tool call → one observation → final normalized date.
- Difference from T1: requires linguistic span identification before arithmetic.

### T3: Constraint-Based Schedule Generation

- Fully iterative, multi-step reasoning loop.
- Planner decides next action:
  - Find start date / Apply intervals / Check rule violations.
- Observation is evaluated by decision module → returns structured state + continue/stop signal.
- Loop continues until constraints are satisfied (or terminates if the 10-turn limit is exceeded).



EWHA WOMANS UNIVERSITY

## IV. Experiments

# Experiment Design: Method Implementation

## Comparison Setup

### Baseline (CoT)

- Few-shot + Chain-of-Thought prompting
- No tool use; all reasoning done internally
- Fixed JSON output schema for all tasks
- One-pass generation (no intermediate verification)

### ReAct Agent (Method-Aligned)

- Implements the **Thought** → **Action** → **Observation** loop
- **T1 / T2:**
  - Single-step reasoning
  - Thought: interpret temporal phrase/sentence
  - Action: one call to *calculator* or *calendar\_db*
  - Observation: tool returns the normalized date → final output
- **T3:**
  - Fully iterative multi-step planning
  - Thought: decide next operation based on constraints & partial schedule
  - Action: tool calls for date arithmetic, holiday lookup, or external event search
  - Observation: decision module evaluates constraint satisfaction and updates state
  - Loop continues until a complete valid schedule is produced
- Modular design separates planning, tool execution, and validation



EWHA WOMANS UNIVERSITY

# Experiment Design: Logistics

## Datasets

- **T1:** 500 Korean relative-time *phrases* (single-date normalization)
- **T2:** 500 Korean relative-time *sentences* (embedded temporal span extraction)
- **T3:** 500 Korean multi-constraint scheduling queries (intervals, weekdays, holidays, counts)

## Models Evaluated

- **GPT-4.1-mini** (via OpenAI API)
- **Solar Pro 2** (via Upstage API)



EWHA WOMANS UNIVERSITY

# Experiment Design: Evaluation Metrics

## Task Accuracy

- **T1 & T2 (Single-Date Tasks)**
  - Output: one ISO-8601 date
  - **Exact-match accuracy:** prediction must match gold date *exactly*
  - No credit for partial matches or  $\pm 1$ -day errors
- **T3 (Schedule Generation)**
  - Output: ordered list of dates
  - **Exact full-sequence match** required
  - Any mistake (wrong date, missing/extra dates, interval errors, constraint violation) → **failure**
  - ReAct exceeding **10-turn limit** also counted as failure

## Efficiency Metrics

- **Latency:** wall-clock time from first prompt to final output
  - Captures extra overhead from multi-step reasoning in T3
- **Token Usage:** total prompt + completion tokens
  - Compares computational cost of CoT vs. ReAct reasoning loops

## Overall Evaluation

- Accuracy = proportion of predictions passing the exact-match criteria
- Latency + token usage provide a measure of **efficiency vs. correctness trade-offs**



## V. Results

# T1 Performance

- **ReAct improves GPT's accuracy:**
  - GPT-ReAct outperformed GPT-CoT by **+4.0 points** (76.2% vs. 72.2%), a **statistically significant gain** ( $p = .0018$ )
  - Shows that ReAct meaningfully reduces single-event reasoning errors in GPT models
- **Solar models show a different trend:**
  - Solar-ReAct achieved **+10 points** over Solar-CoT (76.2% vs. 66.0%) **but without statistical significance** ( $p > .05$ )
  - Suggests Solar benefits from ReAct, but performance variance is higher
- **Efficiency trade-offs differ sharply across models:**
  - **Solar-CoT** = fastest and cheapest (~1.7s latency, low tokens), but lowest accuracy
  - **Solar-ReAct** = best balance (~2.6–2.8s, 2.3–2.4k tokens) with accuracy comparable to GPT-ReAct
  - **GPT-ReAct** = highest accuracy but worst cost efficiency (~25k tokens, ~3.4s)
  - Highlights a **clear accuracy–cost trade-off**: the best-performing model is also the most expensive per query
- **Key insight:**
  - ReAct is effective for single-step tasks, but **its benefit depends heavily on the model backbone**
  - GPT gains accuracy at high cost; Solar gains efficiency with modest accuracy improvements

Model	Accuracy (%)	Avg Latency (s)	Avg Token Usage
GPT-CoT	72.20	~5.11	~2,388
GPT-ReAct	76.20	~3.40s	~25,000
Solar-CoT	66.00	~1.68s	~1,376
Solar-ReAct	76.20	~2.6–2.8s	~2,350–2,450



# T2 Performance

- **GPT-ReAct delivers the strongest accuracy (79%), outperforming GPT-CoT by +4.6 points**
  - Improvement is **statistically significant** (McNemar  $p = .030$ )
- **Solar-CoT shows the lowest accuracy (66.7%) but remains the most computationally efficient**
  - ~1.6s latency, minimal generation (1,376 tokens)
- **Solar-ReAct improves accuracy to 74.85%, matching GPT-CoT-level performance**
  - But **not statistically significant** vs its CoT baseline ( $p > .05$ )
  - Maintains **low cost footprint** (2.5–3.0s latency, ~2.3–2.5k tokens)
- **GPT-ReAct trades extremely high cost for accuracy**
  - ~3.25s latency, **~25k tokens** per sample
  - Lowest **cost-normalized efficiency** despite best accuracy

- **Overall pattern:**
  - ReAct helps GPT meaningfully on T2
  - Solar benefits modestly but not significantly
  - GPT-ReAct gains accuracy through **very long reasoning traces** (tool loops, self-repair, error handling)

Model	Accuracy (%)	Avg Latency (s)	Avg Token Usage
GPT-CoT	74.40	~5.32	~2,393
GPT-ReAct	79.00	~3.25	~25,000
Solar-CoT	66.70	~1.61	~1,376
Solar-ReAct	74.85	~2.5–3.0	~2,350–2,460



# T3 Performance

- All models performed poorly on multi-event temporal reasoning
  - GPT-CoT: **22.2%**
  - GPT-ReAct: **21.4%**  
→ No improvement (McNemar  $p \approx 1.00$ )
- Solar models also struggled
  - Solar-CoT: **16.8%** — lowest overall, but computationally light (~**2.85s** latency, ~**1.3k** tokens)
  - Solar-ReAct: **21.4%**, but not statistically significant vs CoT  
Comes with major cost: **25–30s** latency, **21k–30k** tokens
- GPT-ReAct showed the worst cost–performance trade-off
  - **38–50s** latency
  - **18k+** tokens generated
  - Accuracy still **21.4%** → high cost, no gain

- Overall pattern:
  - ReAct **does not help** in multi-event reasoning
  - **Large computational expansions** (loops, tool failures, self-repair)
  - Models fail to scale beyond single-event tasks

Model	Accuracy (%)	Avg Latency (s)	Avg Token Usage
GPT-CoT	<b>22.20</b>	<b>~1.55</b>	<b>~768</b>
GPT-ReAct	<b>21.40</b>	<b>~38–50</b>	<b>~18,000</b>
Solar-CoT	<b>16.80</b>	<b>~2.85</b>	<b>~1,350</b>
Solar-ReAct	<b>21.40</b>	<b>~25–30</b>	<b>~21,000–30,000</b>



# Error Analysis

## Key Failure Modes

- **CoT Failure Patterns**
  - Frequent ±1-day drift in relative date calculations
  - **Inconsistent Korean week-boundary interpretation** (e.g., Monday vs. “이 번 주/다음 주”)
  - **Reasoning trace vs. final answer mismatch**, indicating unstable internal grounding
  - **Premature reasoning termination**, leading to partially resolved outputs
- **ReAct Failure Patterns**
  - **Tool-driven instability** (bad tool outputs propagate through later turns)
  - **Turn-limit collapses** (hitting 10-step cap before converging)
  - **Redundant re-evaluation loops** and repeated tool calls
  - **Incomplete constraint chaining**, especially in multi-constraint scheduling
- **Task-Level Observation**
  - Problems **amplify sharply in T3**
  - Tool use often **increased**, rather than reduced, error density
  - Indicates that **tool invocation ≠ structured planning** in multi-event reasoning



# Ablation Study: Effect of Tool Flexibility

## Objective

- To investigate whether the performance bottleneck originates from reasoning limitations or tool rigidity.
- Comparing **Hand-coded Tools** (Original) vs. **LLM-driven Tools** (Modified).

## Experimental Results

- **Original Solar-ReAct (Rigid Tools)**
  - Accuracy: **21.40%** (107 / 500)
  - Failure Cases (Max-turn limit exceeded): **218** cases
- **Modified Solar-ReAct (LLM-driven Tools)**
  - Accuracy: **31.40%** (157 / 500)
  - Failure Cases (Max-turn limit exceeded): **85** cases

## Key Findings

- **Performance Gain:** Accuracy improved by +10.0%.
- **Error Reduction:** Significant decrease in "turn-limit collapse" errors (218 → 85).



# Hypotheses Validation Summary

**H1:** “The more complex the constraints, the higher the accuracy and constraint satisfaction in ReAct Agents.”

**Evaluation:** Not supported

**Why:**

- **Single-event tasks (T1/T2):**  
ReAct *did* improve accuracy for GPT (statistically significant) and modestly for Solar (not significant).  
→ This is *positive*, but these tasks had **low complexity**.
- **Multi-event tasks (T3):**  
These are the **most complex constraints**, involving multi-step temporal reasoning.  
Across all models:
  - GPT-CoT: **22.2%**
  - GPT-ReAct: **21.4%** (no improvement,  $p \approx 1.0$ )
  - Solar-CoT: **16.8%**
  - Solar-ReAct: **21.4%** (not significant; massive token cost)
- **ReAct did not increase accuracy under high constraint complexity.**  
In fact, performance **collapsed**, and ReAct often made things *worse* by entering long, unproductive loops.

**Conclusion:**

ReAct does **not** show higher accuracy when constraints grow more complex.

Instead, accuracy **falls sharply**, disproving H1.

**H2:** “ReAct may increase response time and token usage, but shows Pareto-efficient accuracy-to-cost performance.”

**Evaluation:** Partially supported (supported for Solar, contradicted for GPT)

**What the data says:**

**GPT: Not Pareto-efficient**

- Token usage: ~25k per response
- Latency: **3.2–3.4s**
- Accuracy: **79% (T2)**
- **Cost skyrockets**, and accuracy gains are small (+4–5 points).
- In T3, GPT-ReAct has massive cost increases (18k–50k tokens, 38–50s latency) with **zero accuracy benefit**.

→ GPT-ReAct is **not Pareto-efficient**.

**Solar: Largely Pareto-efficient**

- **Solar-CoT:** cheapest but least accurate (66–67%).
- **Solar-ReAct:** moderate cost (2.5–3s, ~2.3–2.5k tokens) with meaningful (but not significant) accuracy gains (74–76%).
- Solar-ReAct essentially matches **GPT-CoT accuracy at a fraction of the cost**.

→ Solar-ReAct appears to be **near-Pareto optimal**: small cost increase → measurable benefit.

**Conclusion:**

- **H2 holds for Solar**, where ReAct produces a good accuracy–cost trade-off.
- **H2 is contradicted for GPT**, where ReAct explodes token usage with diminishing returns.

→ Overall: **partially supported**.



**EWHA WOMANS UNIVERSITY**

## VI. Conclusion

# Summary of Findings

- **Overall Goal:** Compare CoT vs. ReAct for Korean temporal-reasoning tasks.
- **T1 & T2 — Single-event reasoning**
  - ReAct outperforms CoT by **4-6 percentage points**.
  - ReAct shows **lower latency** and **lower token usage**.
  - → **Clear accuracy–cost advantage** for ReAct in simple normalization tasks.
- **T3 — Multi-event, document-level scheduling**
  - Both CoT and ReAct show **low accuracy (17–22%)**.
  - **No significant difference** between the two methods.
  - → Neither approach scales well to multi-constraint schedule generation.
- **Ablation Study**
  - Switching to LLM-driven tools significantly reduced failures, increasing accuracy to **31.4%**.  
→ Suggests that the original low performance stemmed from "**Tool Misuse**" due to **rigid specifications**, rather than a lack of reasoning power.



# Implications

- Evaluation of reasoning frameworks must account for **accuracy, cost-efficiency, and alignment with task structure**, not just raw correctness.
- **ReAct is effective for lightweight, single-event reasoning (T1/T2)** — delivers strong accuracy gains with *lower* latency and token cost.
- **ReAct is not sufficient for multi-event temporal reasoning (T3)** — tasks require **explicit event segmentation, constraint propagation, and hierarchical planning**, which ReAct does not provide.
- **Ablation findings show tool usage is not always beneficial** — performance improved *without* tool invocation, indicating that **stable temporal reasoning depends more on orchestration and planning** than on adding external tools.
- Future systems should prioritize **structured controllers, planning modules, and temporal decomposition pipelines** for complex Korean scheduling tasks.



# Limitations & Future Work

- Current evaluation focuses primarily on **accuracy**, leaving deeper **qualitative analysis** of reasoning traces and failure patterns for future study.
- Task difficulty varies: **T1/T2 contain simpler temporal structures**, which may partly explain the performance gap compared to T3.
- T3 exposes the need for richer temporal modeling; future work should explore:
  - **Multi-event pipeline architectures** for complex temporal queries
  - **Adaptive tool-use policies and error-recovery mechanisms**
  - New metrics capturing **temporal reasoning stability** beyond exact match
  - Broader coverage of **Korean temporal expressions**, including discourse-level and culturally specific time references



Thank you.

# References

- [1] Cho, M. J., Namgung, S., & Kim, M. Y. (2022). A Rule-based Approach for Recognition and Normalization of Temporal Information in Korean Text. *Journal of the Korea Institute of Information and Communication Engineering*, 26(11), 1539-1547.
- [2] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). *ReAct: Synergizing reasoning and acting in language models*. In The Eleventh International Conference on Learning Representations (ICLR). arXiv. <https://arxiv.org/abs/2210.03629>
- [3] Jeong, Y.-S., Kim, Z. M., Do, H.-W., Lim, C.-G., & Choi, H.-J. (2015). *Temporal information extraction from Korean texts*. In Proceedings of the 19th Conference on Computational Natural Language Learning (CoNLL 2015) (pp. 250–260). Association for Computational Linguistics. <https://aclanthology.org/K15-1028.pdf>
- [4] Piryani, B., Abdallah, A., Mozafari, J., Anand, A., & Jatowt, A. (2025). *It's high time: A survey of temporal information retrieval and question answering*. arXiv. <https://arxiv.org/abs/2505.20243>
- [5] Eldan, O., & Yahav, A. (2025). *LM-Polygraph: Unveiling Instability in Large Language Model Leaderboards*. arXiv preprint arXiv:2506.21595.
- [6] Cherukuri, M. (2025). Cost, Complexity, and Efficacy of Prompt Engineering Techniques for Large Language Models. *International Journal on Science and Technology (IJSAT)*, 16(2).
- [7] IBM. (2024). *What is a ReAct agent?*. Retrieved October 21, 2025, from <https://www.ibm.com/think/topics/react-agent>
- [8] Agent Patterns Documentation. (2024). *ReAct Pattern*. Retrieved October 21, 2025, from [https://agent-patterns.readthedocs.io/en/stable/patterns/re\\_act.html](https://agent-patterns.readthedocs.io/en/stable/patterns/re_act.html)
- [9] Lim, C. G., Jeong, Y. S., & Choi, H. J. (2018). *Korean TimeBank Including Relative Temporal Information*. In Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018).



## Appendix: Related Works

# Temporal Information Extraction from Korean Texts

(Jeong, Y.-S., Kim, Z. M., Do, H.-W., Lim, C.-G., & Choi, H.-J.; 2015) [3]

## Why Korean Temporal Expressions Are Challenging

- Korean's linguistic structure complicates normalization:
  - Verb morphology with stems + endings (어 간/어 피)
  - Frequent omission of subjects/objects
  - Multiple numeral systems (삼십 / 서른)
  - Use of lunar-calendar expressions
- These characteristics increase ambiguity for LLMs and rule-based systems.
- Accurate normalization is critical for age/interval questions and context-dependent references used by AI agents.

## Main Contributions of K15-1028

- One of the earliest systems dedicated to **Korean** temporal expression recognition.
- Proposes a **two-stage pipeline**:
  - Segment text into minimal semantic units.
  - Apply **finite-state patterns** to identify temporal expressions.
- Achieves high precision/recall even in noisy, domain-specific Korean text.
- Demonstrates feasibility of structured detection despite Korean-specific linguistic challenges.

## Limitations / Gap

- Rule-based method struggles with conversational or informal Korean.
- Doesn't address deep temporal reasoning or multi-step interval logic.
- Leaves open the need for **LLM-based** Korean temporal normalization.



EWHAWOMANS UNIVERSITY

# It's High Time: A Survey of Temporal QA

(Piryani, B., Abdallah, A., Mozafari, J., Anand, A., & Jatowt, A.; 2025) [4]

## Why Korean Temporal Reasoning Needs Attention

- Temporal QA is complex even in English; even more so in Korean due to:
  - Morphological variation and omitted arguments
  - Multiple numeral systems
  - Relative or contextual expressions (“지난번”)
  - Lunar-calendar dates
- LLMs frequently misinterpret these forms, limiting accurate time anchoring for Korean tasks.

## Main Points From the Survey

- Identifies universal challenges in Temporal QA:
  - Normalizing relative expressions
  - Anchoring events to timelines
  - Reasoning over intervals and sequences
- Notes that even advanced LLMs struggle with temporal consistency and multi-step logic.
- Rates current systems as inadequate for languages with complex temporal morphology.

## Relevance to Our Project

- Survey highlights the need for **structured reasoning** frameworks to improve temporal QA.
- Korean-specific challenges amplify this need, as ambiguity and morphology make simple CoT unreliable.
- Motivates exploring **ReAct-style** reasoning loops for Korean temporal normalization and reasoning.



Ewha Womans University

# ReAct: Synergizing Reasoning and Acting

(Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y; 2023) [2]

- Integrates **Reasoning + Acting** through an iterative loop:  
**Thought → Action → Observation**
- Overcomes limitations of purely CoT or purely action-based methods
  - Reduces hallucinations
  - Adds intermediate verification
  - Enables adaptive decision-making

## Advantages Shown in Prior Research

- Improved performance on:
  - Multi-hop QA
  - Fact verification
  - Interactive reasoning tasks
- Offers better factual grounding, interpretability, and self-correction.

## Relevance to Korean Temporal Reasoning

- Temporal reasoning often requires:
  - Constraint checking
  - Step-by-step validation (e.g., calendars, dates, offsets)
  - Handling ambiguous or context-dependent expressions
- ReAct's structured loop may enhance the accuracy and robustness of Korean temporal normalization.



Ewha Womans University